



Cloud Computing

Hardware virtualization

Seyyed Ahmad Javadi

sajavadi@aut.ac.ir

Spring 2023



Introduction

Hardware-level Virtualization

- **An abstract execution environment in terms of computer hardware on top of which a *guest operating system* can be run.**

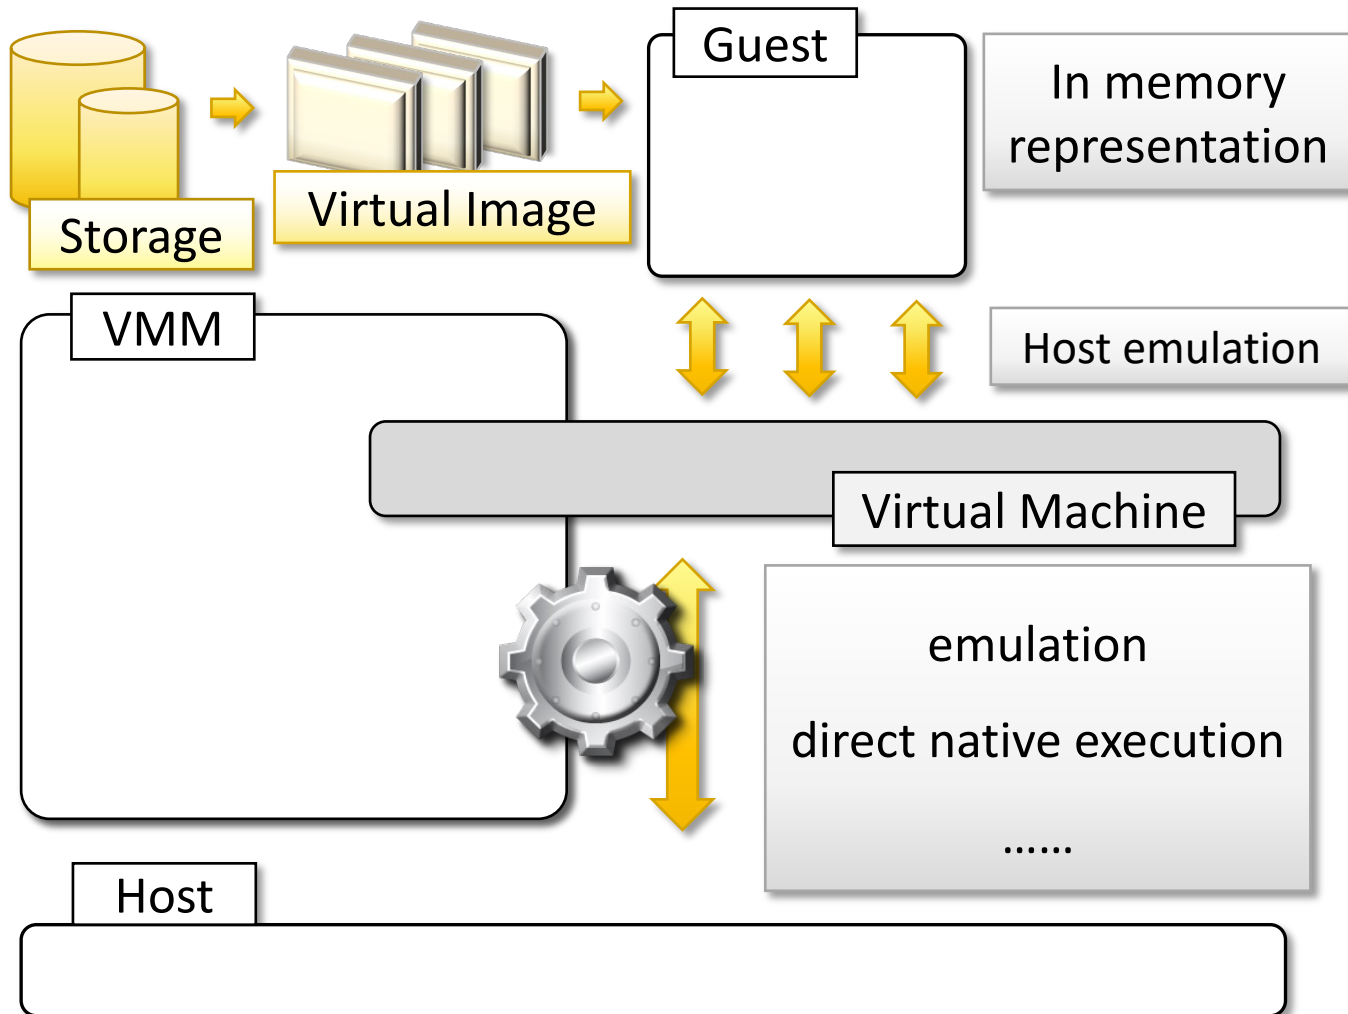
Concept	Represented by
<i>Guest</i>	Operating system
<i>Host</i>	Physical computer hardware
<i>Virtual machine</i>	Its emulation
<i>Virtual machine manager</i>	Hypervisor

What is Hypervisor?

Hypervisor is a program enabling the abstraction of the underlying physical hardware.

Hypervisor is also called Virtual Machine Manager (**VMM**)


Hardware-level Virtualization



Types of Hypervisor

➤ **Type I hypervisors** (native VM)

- Run *directly* on top of the hardware.
- *Take the place* of the operating systems
- Interact directly with the ISA interface



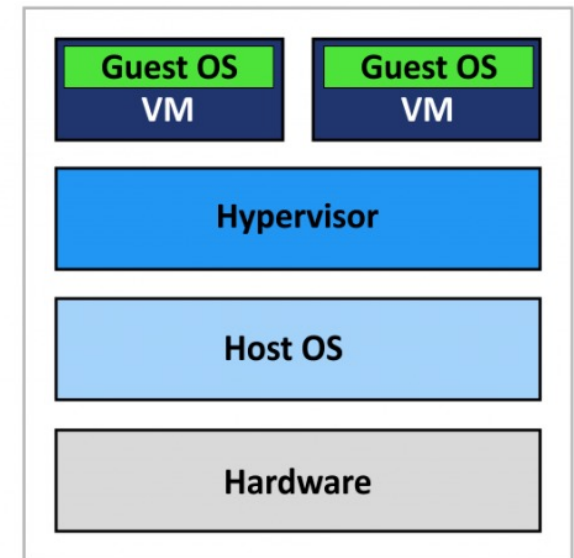
**Type 1 Hypervisor
(Bare-Metal Architecture)**

Source: [http://:
https://www.nakivo.com/blog/hyper-v-
virtualbox-one-choose-infrastructure/](http://:https://www.nakivo.com/blog/hyper-v-virtualbox-one-choose-infrastructure/)

Types of Hypervisor (cont.)

➤ *Type II hypervisors* (hosted VM)

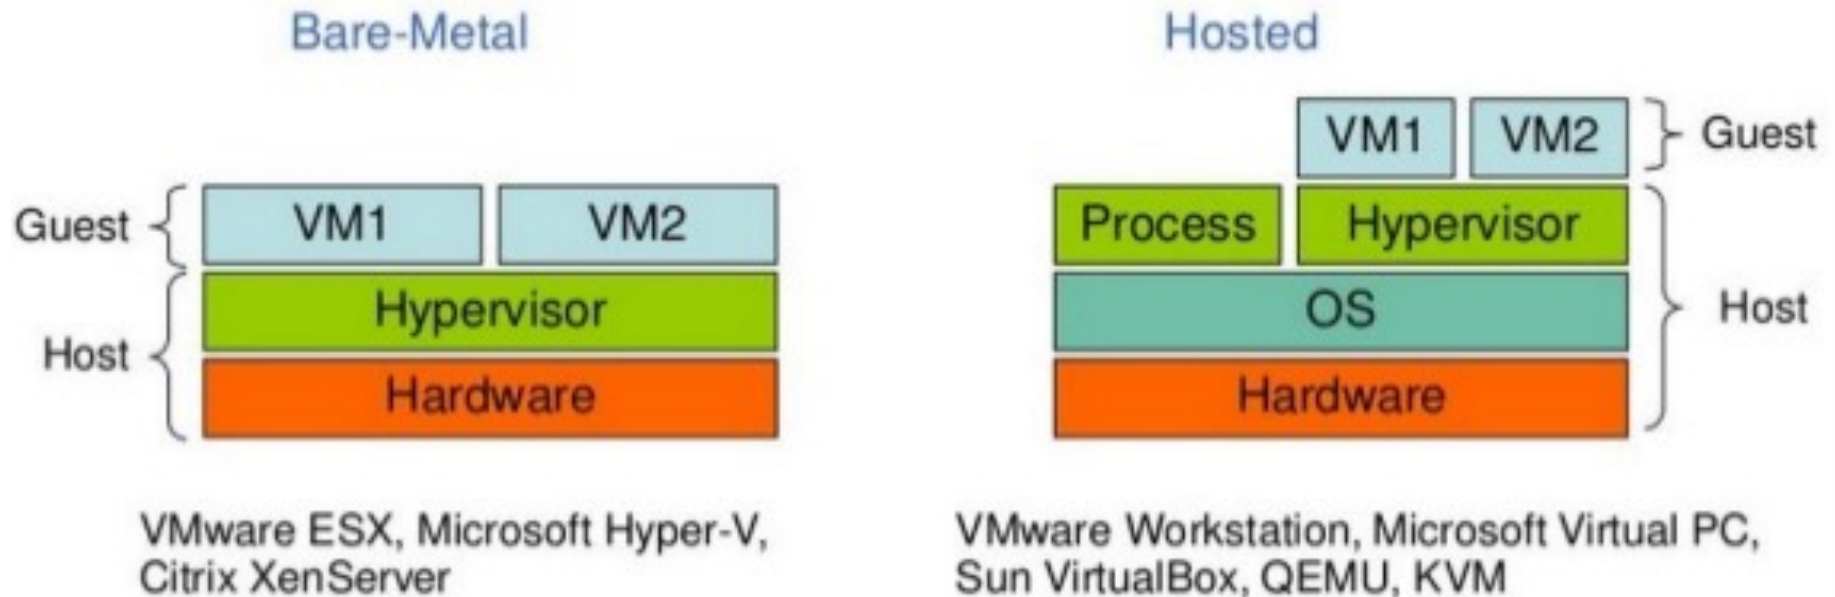
- Require the support of an operating system
- Are programs **managed by the operating system**
- Interact with operating system through the **ABI**.



**Type 2 Hypervisor
(Hosted Architecture)**

Source: [http://:
https://www.nakivo.com/blog/hyper-v-virtualbox-one-choose-infrastructure/](http://https://www.nakivo.com/blog/hyper-v-virtualbox-one-choose-infrastructure/)

Type of Hypervisors (cont.)



Source: <https://www.slideshare.net/PraveenHanchinal/virtualizationthe-cloud-enabler-by-inspiregroups/18-Types of hypervisors VMM>

Approaches of Executing Guest Instructions

Executing Guest Instructions

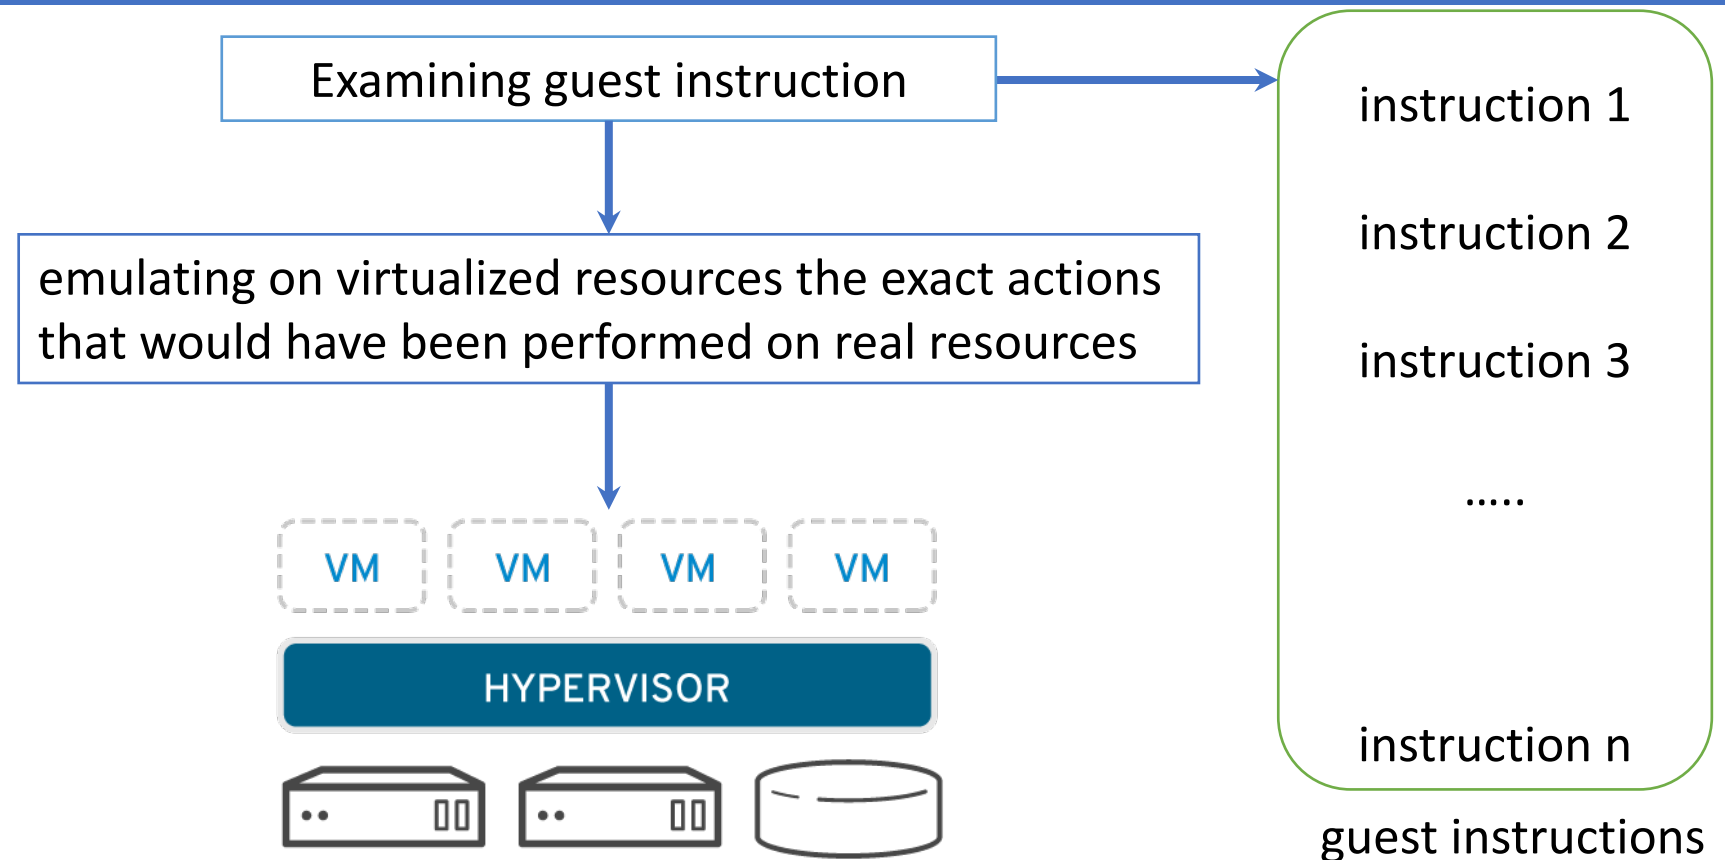
➤ **Emulation**

➤ **Direct native execution**

Emulation

“the process of implementing the interface and functionality of one system or subsystem on a system or subsystem having a different interface and functionality...”

Emulation (cont.)

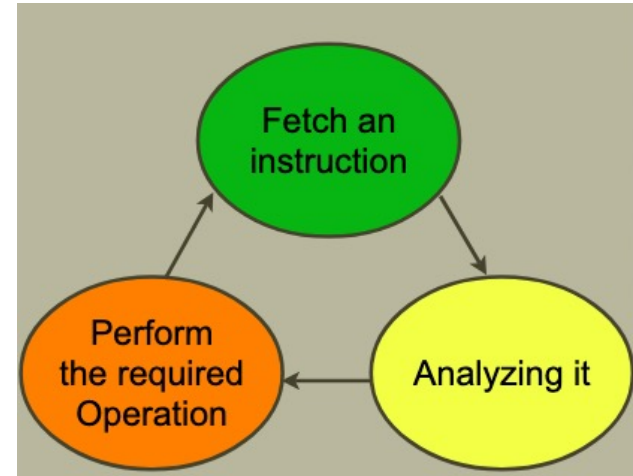


Only available mechanism when the ISA of the guest is *different* from the ISA of the host.

Emulation Approaches

➤ Interpretation

- Done in software,
one instruction at a time




<http://cse.unl.edu/~witty/class/embedded/material/note/emulation.pdf>

➤ Binary translation

- Translating a block of source instructions to target instructions.
- Saving the translated code for repeated use

Interpretation versus Binary Translation



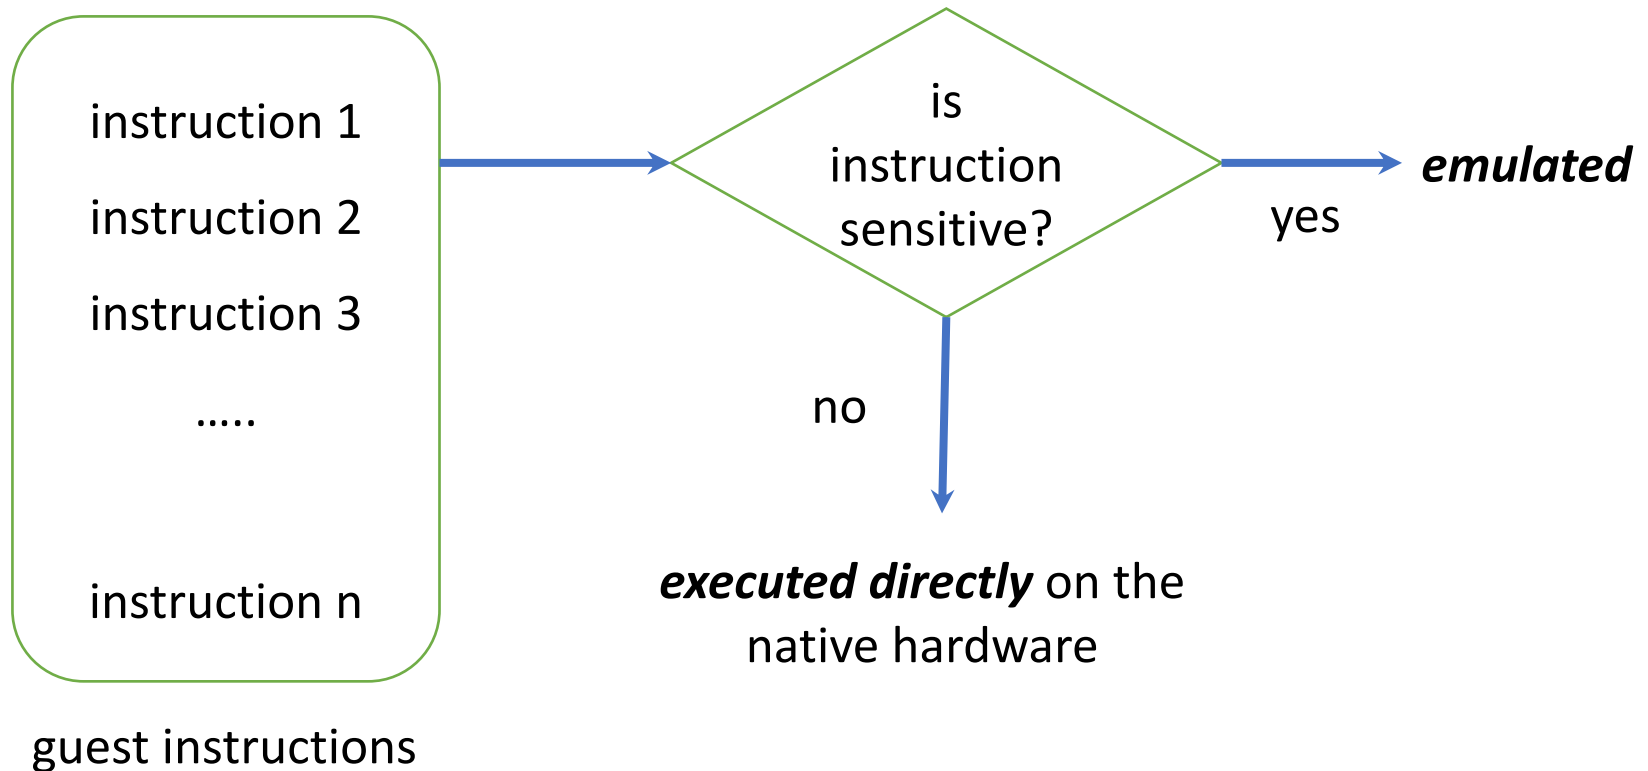
<http://se.unl.edu/~witty/class/embedded/material/note/emulation.pdf>

Interpretation versus Binary Translation (cont.)

	Implementation	Performance
Interpretation	simple and easy	low
Binary Translation	complex	high initial translation cost, small execution cost

<http://www.ittc.ku.edu/~kulkarni/teaching/EECS768/slides/chapter2.pdf>

Direct Native Execution



Only if the ***ISA of the host is identical to the ISA of the guest.***

Hardware Virtualization Methods

Hardware Virtualization Methods

➤ **Full Virtualization**

➤ **Paravirtualization**

Full Virtualization

- Run a program **directly on top of a VM** and **without any modification**
 - The program thought it were run on the raw hardware.
- The principal advantage of full virtualization
 - Complete isolation → enhanced security
 - Ease of emulation of different architectures
 - Coexistence of different systems on the same platform.

Full Virtualization (cont.)

➤ In some architectures, **some sensitive instructions are not privileged**

- They cannot be virtualized in the classic way.
- Like the non-hardware-assisted x86

➤ ***Two technologies:***

- Emulation (Binary Translation and Interpretation)
- Hardware-assisted virtualization


First scenario:
Host ISA \neq Guest ISA

Main Challenge

- Guest Instruction can not directly executed on the host
- Feasible virtualization technologies
 - Interpretation
 - Binary Translation

✓

Binary Translation



Binary Translation (cont.)

➤ Static Binary Translation

- On a full program

➤ Dynamic Binary Translation

- Introduces an additional overhead.

Dynamic Binary Translation

- It is usually performed in small units called "***basic blocks***".
- A basic block is a set of instructions that ends with a branch instruction but does not have any branch instructions inside.
 - Be executed start to finish by a CPU
 - An ideal unit for translation
- The translations of the basic blocks are ***cached***
 - Overhead of translating only happens the first time a block is executed.

<https://blogs.oracle.com/ravello/nested-virtualization-with-binary-translation>

Static vs Binary Translation

	Input type	Granularity	Translation time
Static binary translation	Binary program	Full program	Before running program
Dynamic binary translation	Binary program	Basic block	At runtime

Demo

- Host: Mac M1 chip (ARM-based processor)
- Guest: Ubuntu ISO compiled for X86 processor (AMD)
- Start time of Ubuntu AMD ISO installation: 8:45
- Finish time: 10:05
- Why it takes so long?

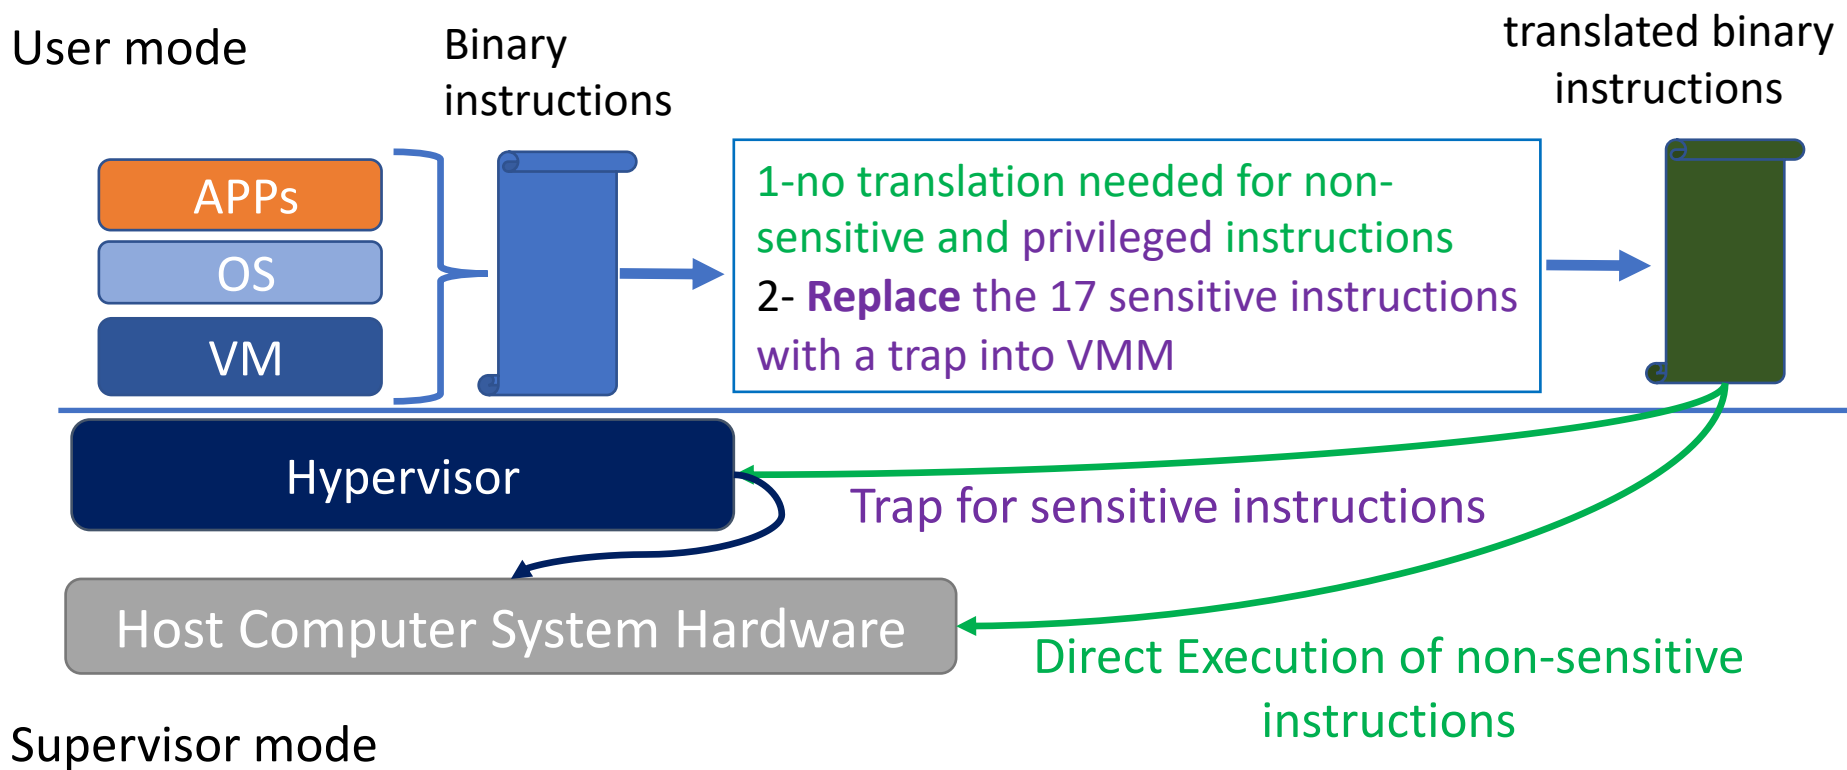
Second scenario:
Host ISA = Guest ISA =
Original ISA Or
Non-hardware assisted X86

Main Challenge

- **Some (17) sensitive** instructions are not privileged.
 - See previous lecture
- They cannot be virtualized in the classic way.
- Feasible virtualization technologies
 - Interpretation
 - Binary Translation and direct execution-based technique

Binary Translation and direct execution-based technique

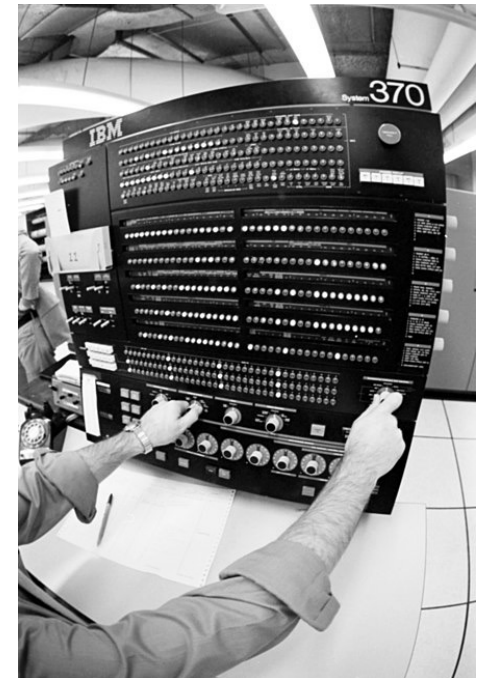
- Replaces the sensitive instructions that do not generate traps (critical instructions) with a trap into the VMM to be emulated in software.



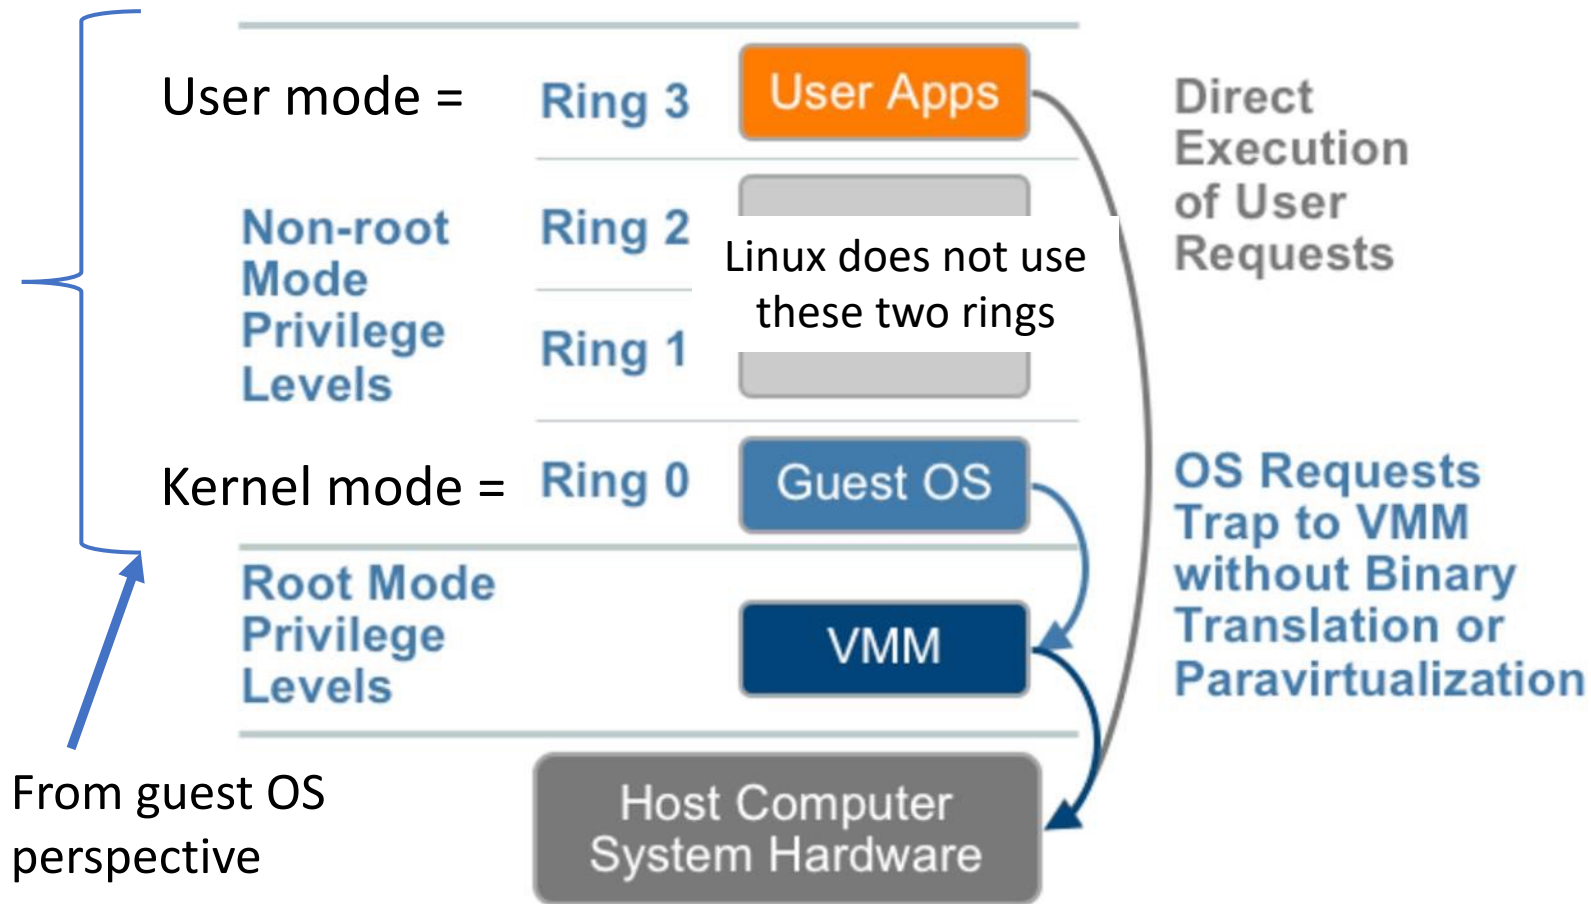
Third Scenario:
Host ISA = Guest ISA =
Hardware assisted X86
or
Hardware assisted ARM

Hardware-assisted Virtualization

- **Architectural support** for building a VMM able to run a guest operating system in complete isolation.
- This technique was originally introduced in the IBM System/370.
- Extensions to x86-64 architecture
 - Introduced with Intel-VT and AMD-V.



Hardware-assisted Virtualization (cont.)



<https://thecustomizewindows.com/2014/09/hardware-assisted-virtualization/>

Intel-VT and AMD-V

- New CPU execution mode feature
- This allows the VMM to run in a new root mode below ring 0
 - **Ring 0P**: privileged root mode (VMM)
 - **Ring 0D** : de-privileged non-root mode (Guest OS)
- Sensitive calls are set ***to automatically trap*** to the hypervisor and handled by hardware
 - Removing the need for either binary translation or para-virtualization.

Intel-VT

- Main feature: inclusion of the new VMX mode of operation.

	all four IA-32 privilege levels (rings)	VMX instructions
VMX non-root operation		
VMX root operation		

VMX Instructions

➤ "VMX" stands for Virtual Machine Extensions

13 new instructions

VMPTRLD	VMPTRST	VMCLEAR	VMREAD	VMWRITE
VMCALL	VMLAUNCH	VMRESUME	VMXOFF	VMXON
INVEPT	INVVPID	VMFUNC		

➤ Permit entering and exiting a ***virtual execution mode*** where the ***guest OS perceives*** itself as running with full privilege (ring 0), but the ***host OS remains protected***.

Hardware-assisted Virtualization

- The behavior of the processor in ***non-root operation is limited*** in some respects from its behavior on a normal processor.
- ***Critical shared resources are kept under the control of a monitor running in VMX root operation.***
- VMM is run in VMX root mode
- Virtual machine and the guest OS are run in non-root mode.

Examples of Hardware-assisted Virtualization

➤ VirtualBox

➤ VMware

➤ Microsoft Hyper-V



Demo

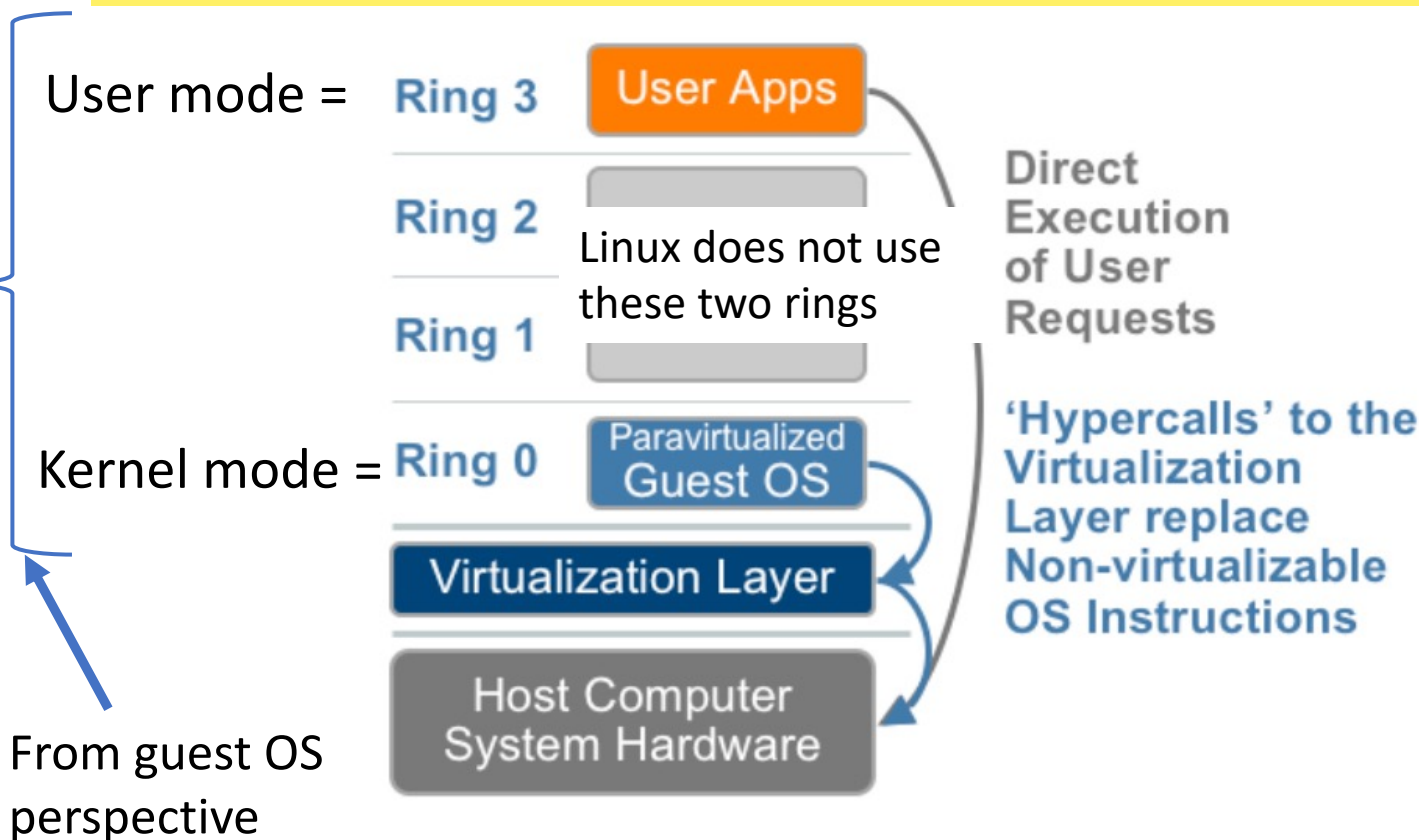
- Host: Mac M1 chip (ARM-based processor)
- Guest: Ubuntu ISO compiled for ARM processor
- Time taken to install Ubuntu ARM ISO: less than 10 mins
- Let's do a simple benchmark

Paravirtualization



Paravirtualization

- Paravirtualization refers to communication between the guest OS and the hypervisor to improve performance and efficiency.



Paravirtualization (cont.)

- It is not a transparent virtualization solution
 - Allows implementing *thin* virtual machine managers.
 - Remapping the performance-critical operations through the virtual machine software interface.

- Expose a software interface to the virtual machine that is slightly modified from the host
 - As consequence, guests need to be modified.

Paravirtualization (cont.)

- Provide the capability to demand the execution of performance-critical operations ***directly on the host***
 - Preventing performance losses that would otherwise be experienced in managed execution.
- Allows a **simpler implementation of virtual machine managers**
 - VMM have to simply **transfer** the execution of performance-critical operations **directly to the host**.
 - These instructions were ***hard to virtualize***

Paravirtualization (Cont.)

- Xen is ***the most popular implementation*** of paravirtualization.



- The guest operating systems need to be changed
- The sensitive system calls need to be re-implemented with ***hypercalls***
 - Are specific calls exposed by the virtual machine interface of Xen.

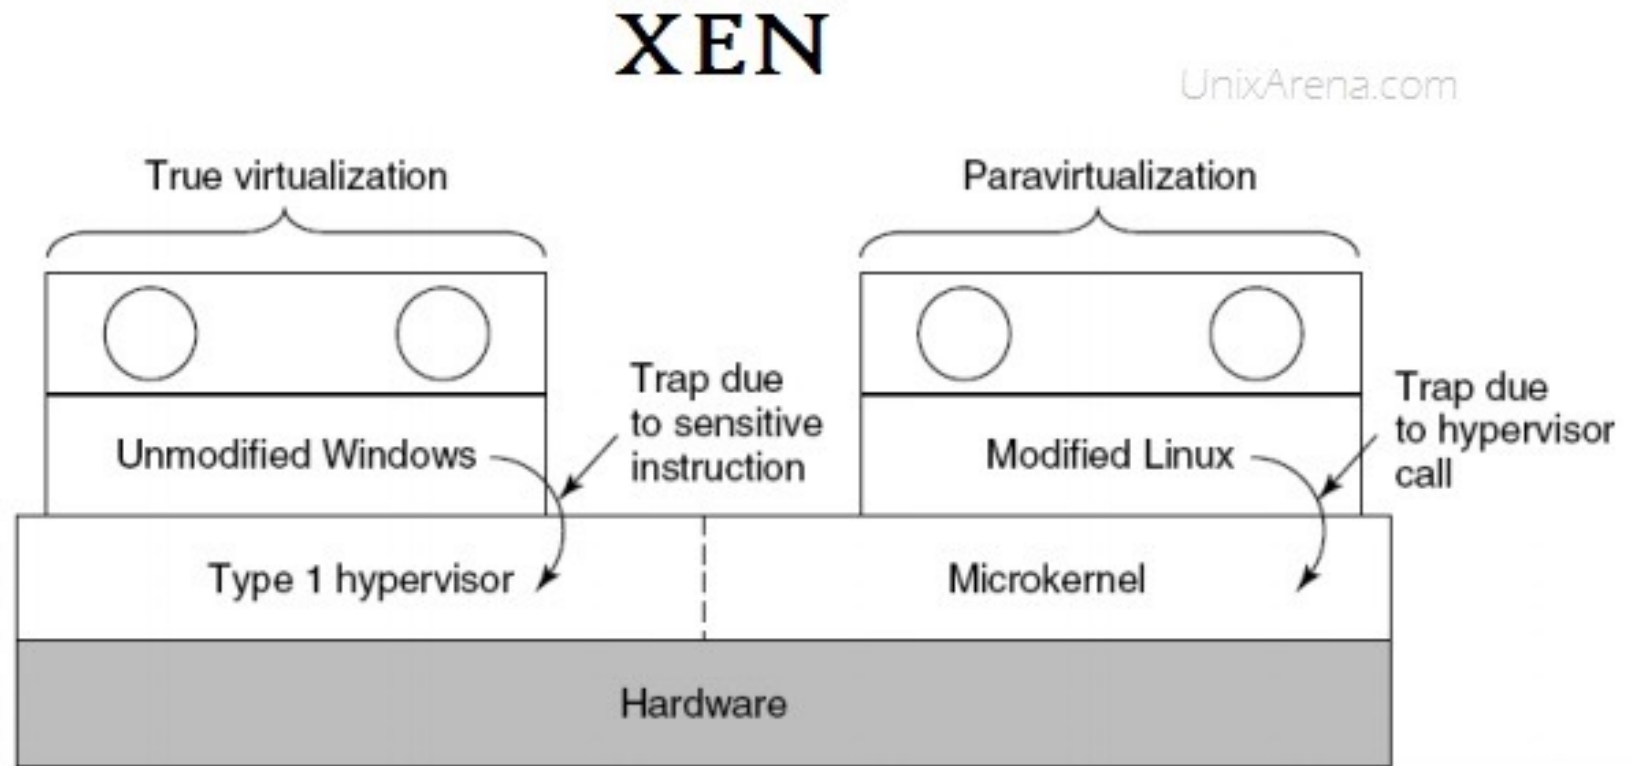
Paravirtualization (Cont.)

➤ With the use of ***hypercalls***, the Xen hypervisor is able to

- **catch the execution of all the sensitive instructions**
- **manage them,**
- **and return the control**

to the guest operating system by means of a supplied handler.

Xen Hypervisor



Xen supports both Full virtualization and Para-virtualization

[source:https://www.unixarena.com/2017/12/para-virtualization-full-virtualization-hardware-assisted-virtualization.html/](https://www.unixarena.com/2017/12/para-virtualization-full-virtualization-hardware-assisted-virtualization.html/)

Paravirtualization (cont.)

- Open-source operating systems such as Linux can be easily modified
 - Their code is publicly available
 - Xen provides full support for their virtualization

- Components of the Windows family ***are generally not supported*** by Xen unless hardware-assisted virtualization is available.