



# **Cloud Computing**

## **OpenStack Nova Architecture**

Seyyed Ahmad Javadi  
[sajavadi@aut.ac.ir](mailto:sajavadi@aut.ac.ir)

Fall 2023

<https://www.slideshare.net/HaimAteya/an-intrudction-to-openstack-2017>

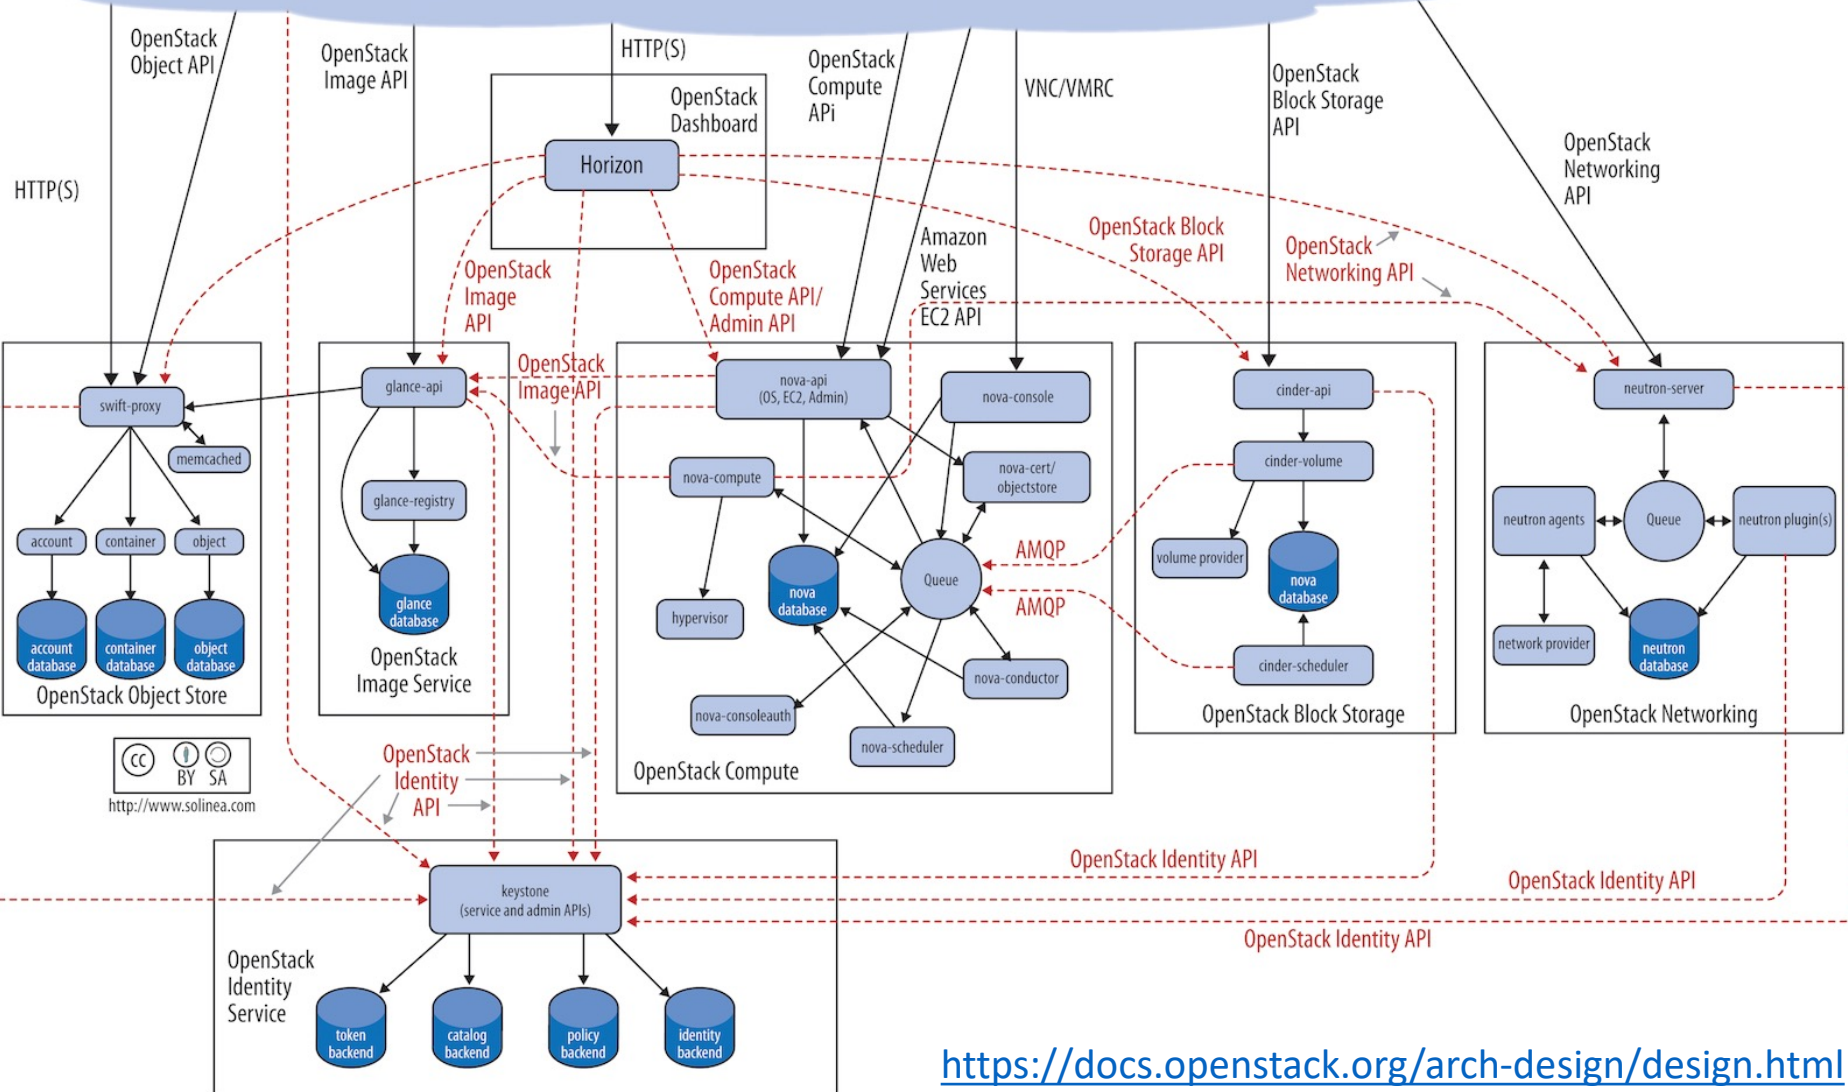
<https://docs.openstack.org/security-guide/introduction/introduction-to-openstack.html>

# Overview



- Command-line interfaces (nova, neutron, swift, etc)
- Cloud Management Tools (Rightscale, Enstratus, etc)
- GUI tools (Dashboard, Cyberduck, iPhone client, etc)

Internet



<https://docs.openstack.org/arch-design/design.html>

# Controller node components

---

glance_api			
horizon			
heat_engine	heat_api_cfn	heat_api	
neutron_server			
ovn_northd	ovn_sb_db	ovn_nb_db	
nova_api	nova_scheduler	nova_novncproxy	nova_conductor
cinder_scheduler	cinder_api		
keystone	keystone_fernet	keystone_ssh	
placement_api	etcd	rabbitmq	memcached
keepalived	haproxy	cron	kolla_toolbox
chrony			

No need to memorize the table for the exam

# Compute node components

---

openvswitch_vswitchd	openvswitch_db	
nova_compute	nova_libvirt	nova_ssh
cinder_backup	cinder_volume	
tgtd	iscsid	
neutron_ovn_metadata_agent	ovn_controller	
cron	kolla_toolbox	chrony

No need to memorize the table for the exam

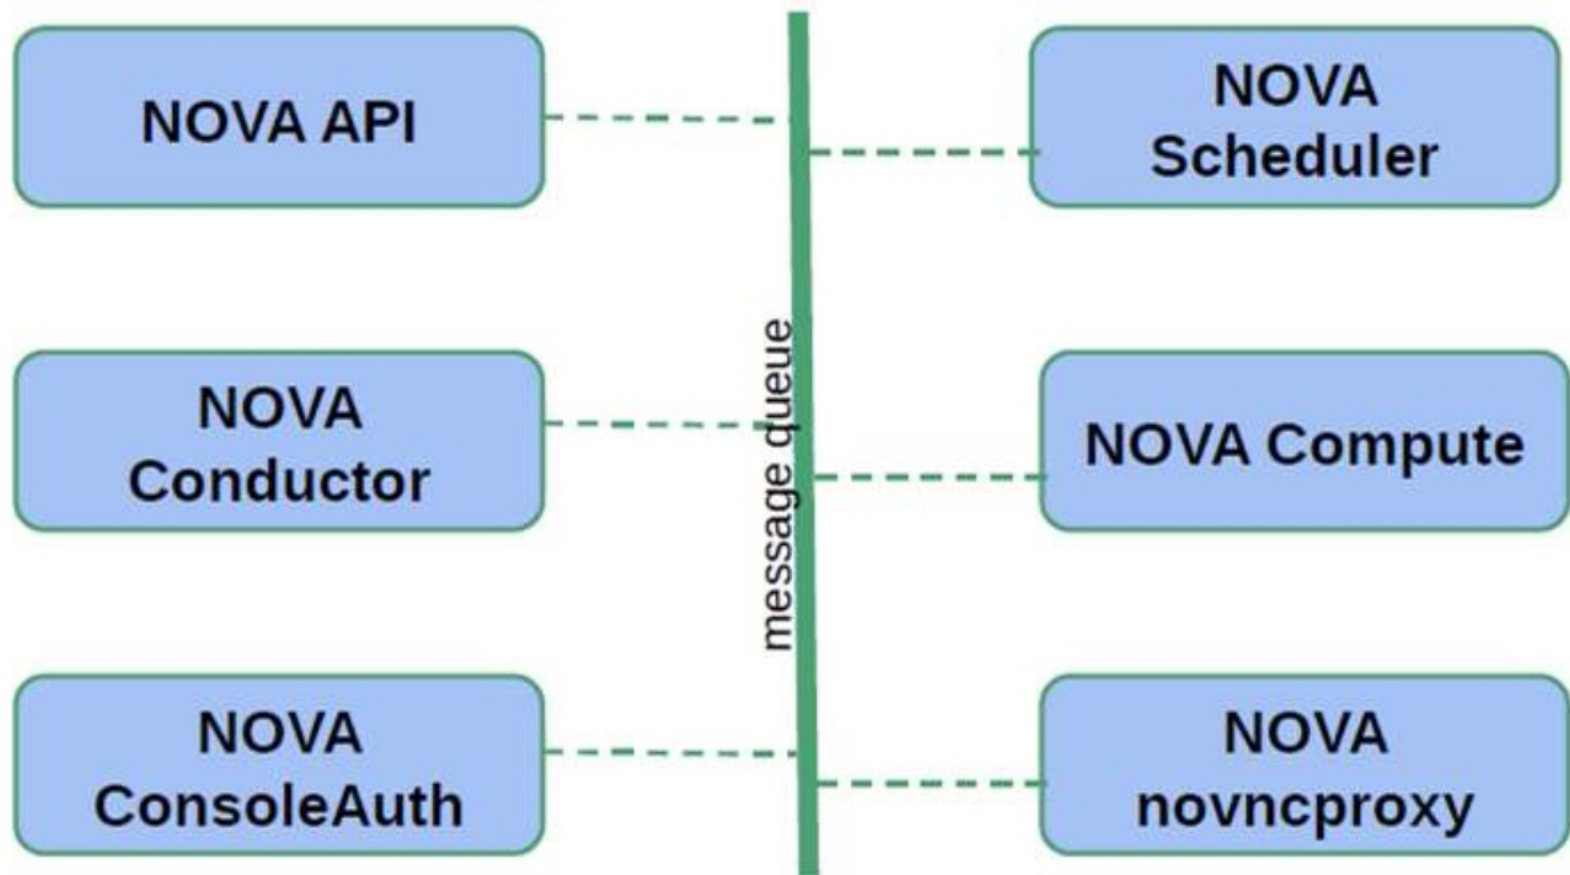
# Nova

---

- Provided compute as service
- The main part of an IaaS system
- It is designed to manage and automate pools of computer resources
- Compute's architecture is designed to scale horizontally

# Nova Components

---



# Nova Components

---

## ➤ Nova-conductor

- Provides database-access support for Compute nodes

## ➤ Nova-consoleauth

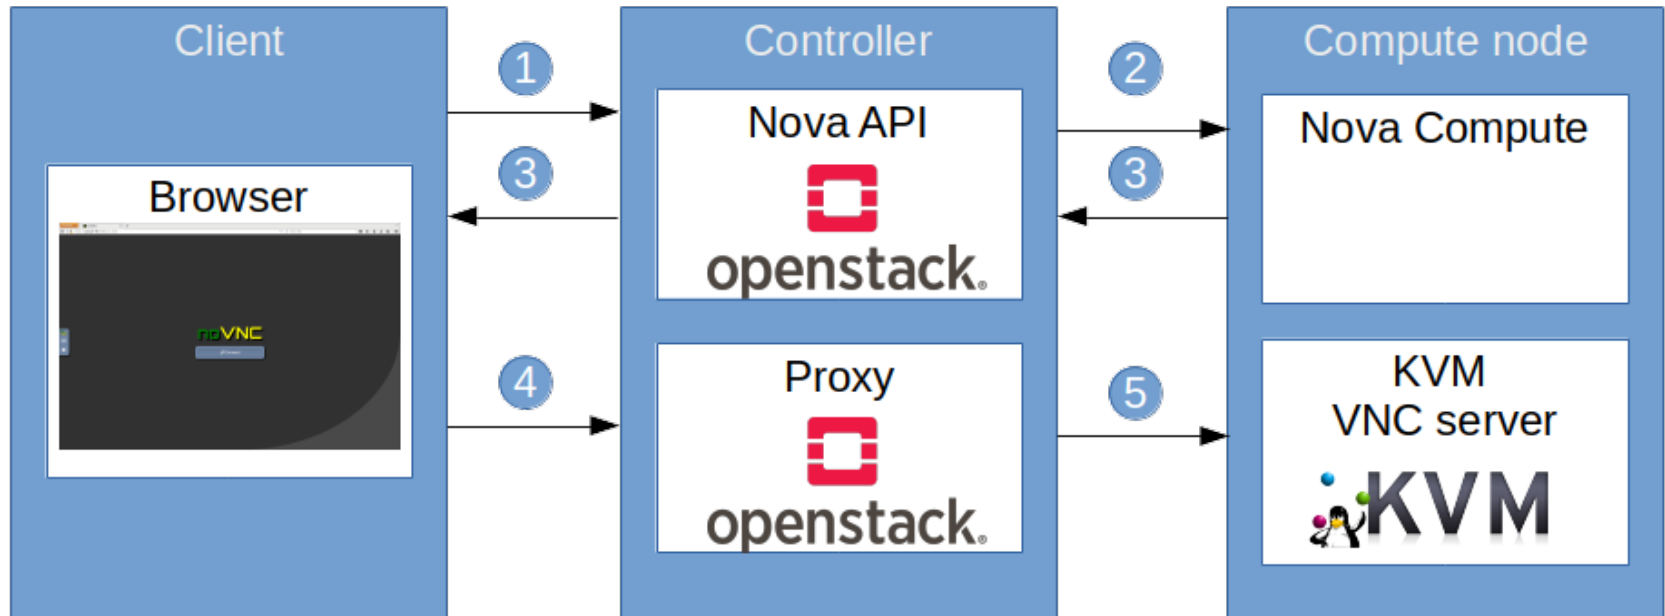
- Handles console authentication

## ➤ Nova-novncproxy

- Provides a VNC proxy for browsers



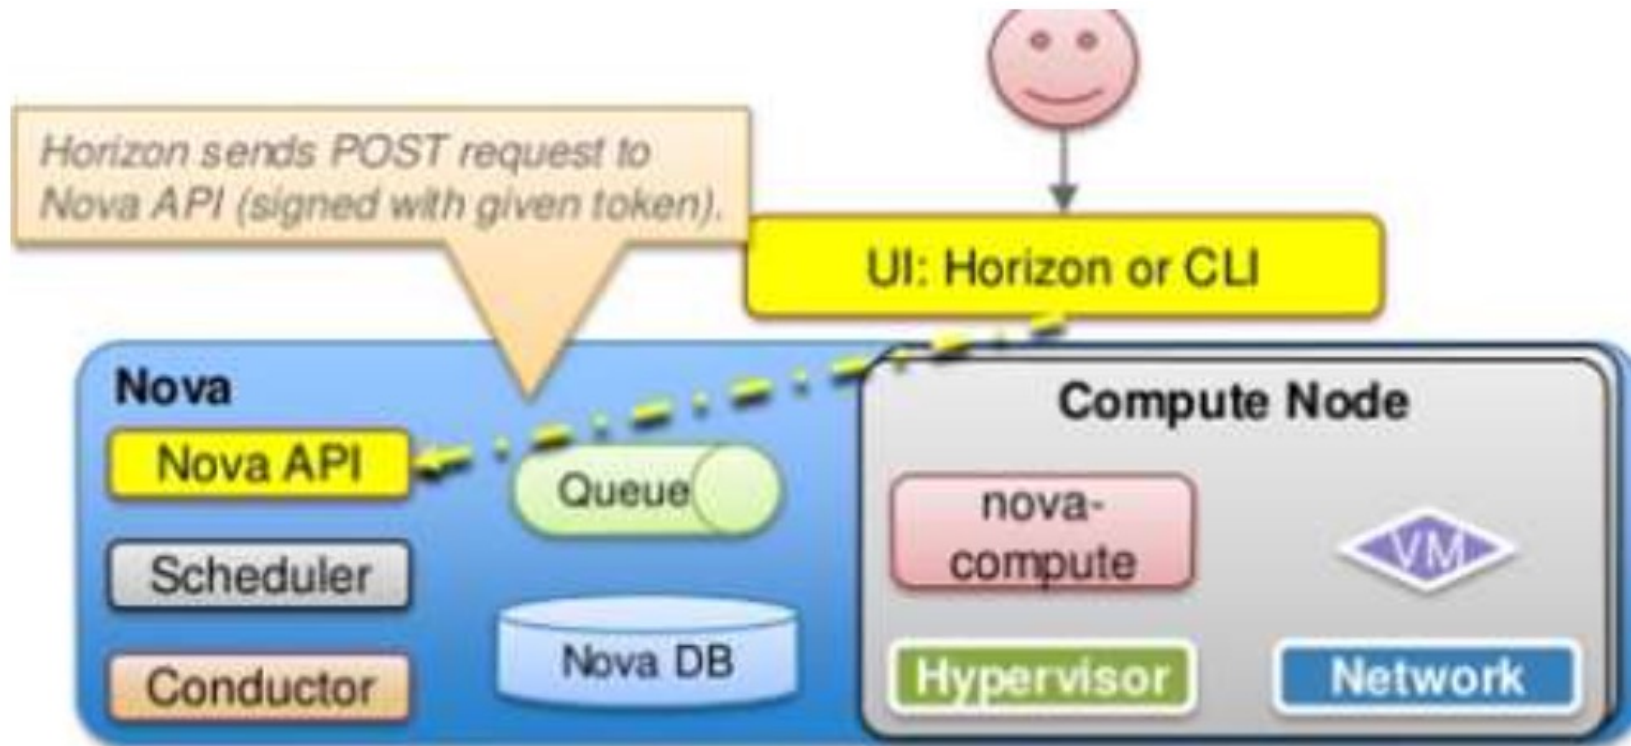
# The Nova VNC proxy



<https://leftasexercise.com/2020/02/14/openstack-nova-installation-and-overview/>

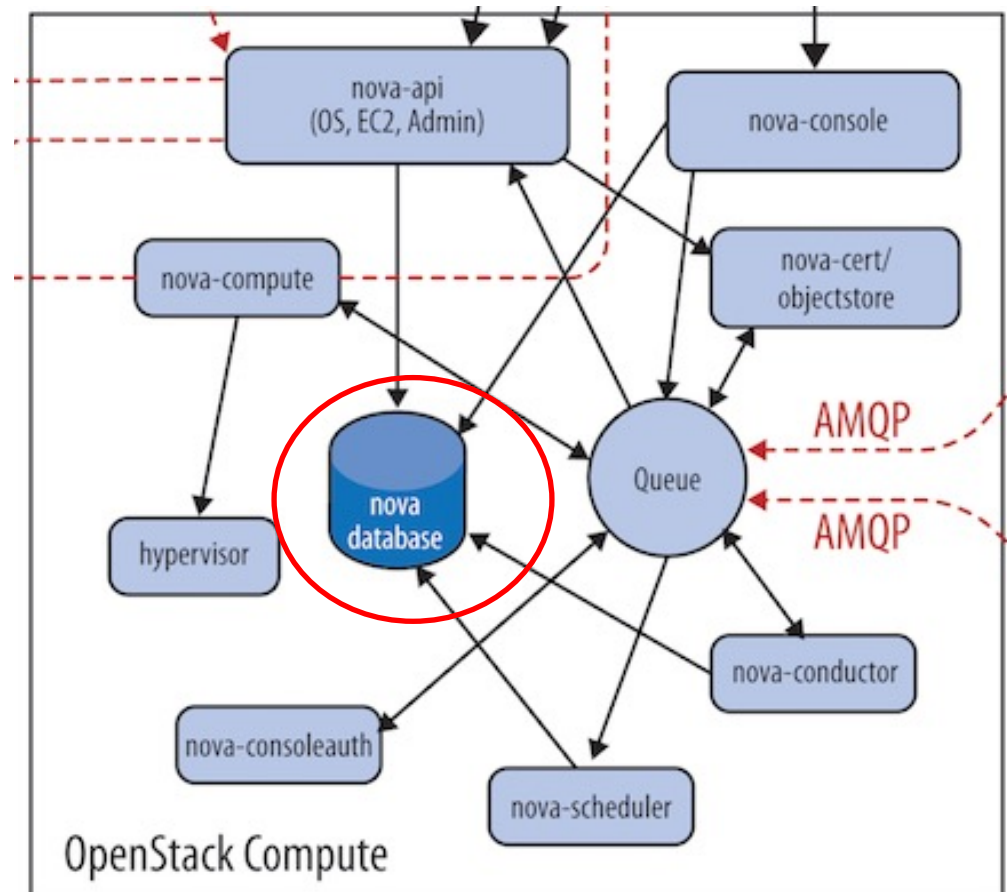
# NOVA API

- NOVA-API is responsible to provide an API for users and services to interact with NOVA



# On Compute Node

- There is a periodic task (Resource Tracker), which collects host information.
- This information is then stored to ***database***

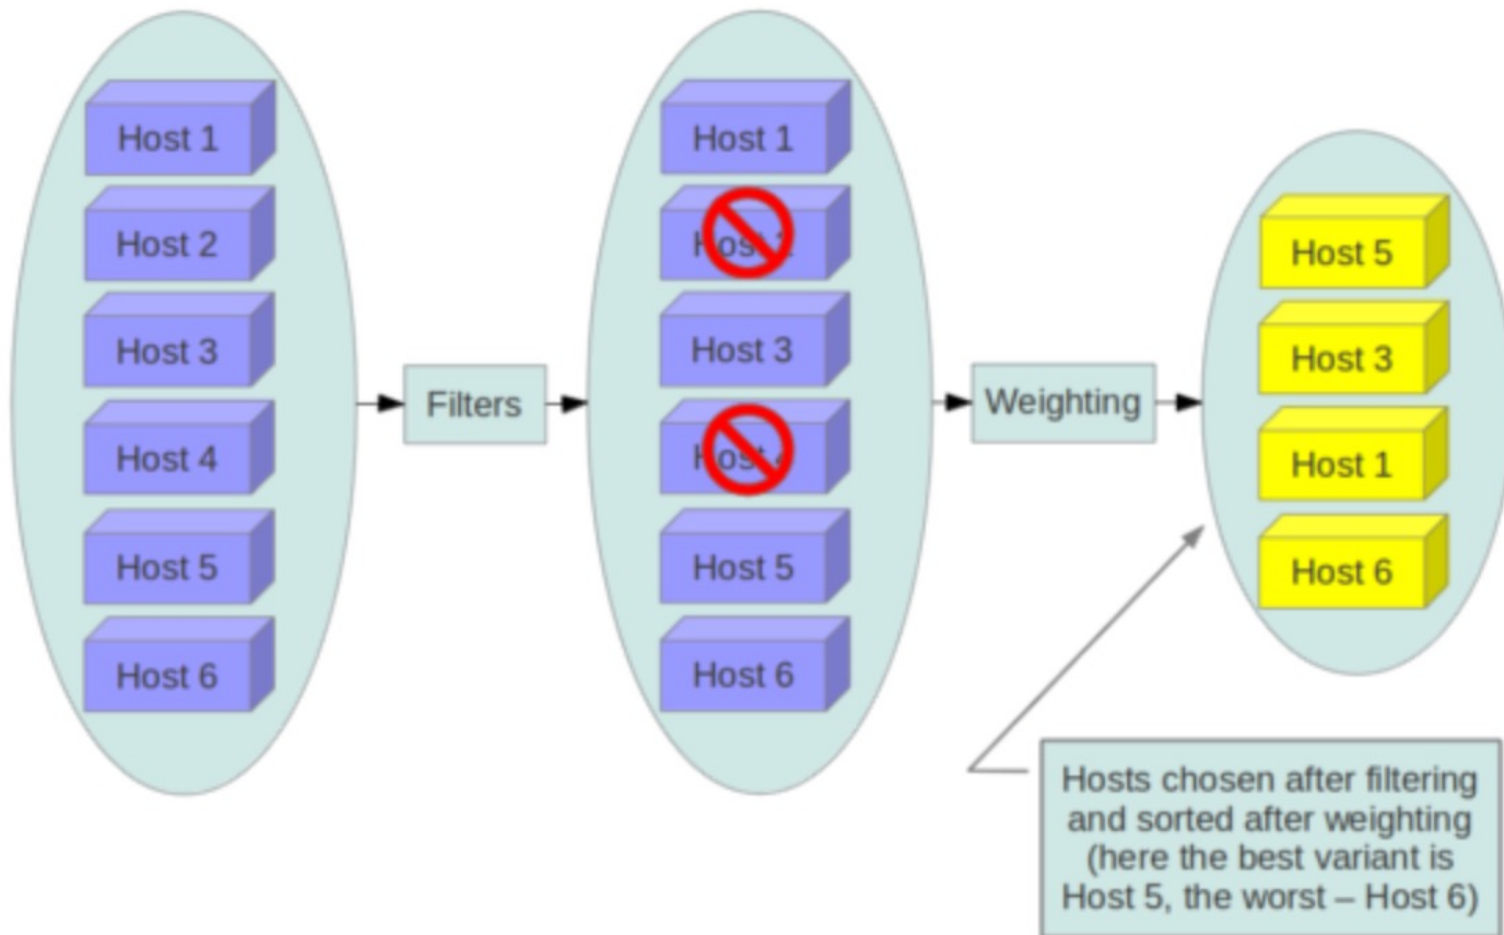


# On Controller Node

---

- Request from nova API reaches conductor
- Conductor interacts with the scheduler
- Scheduler ***uses filters*** to identify the best node
  - From the information stored in ***database***
- Selected host information is sent back to conductor
- Conductor uses the compute queue and directs it to selected host
- The compute node then launches the instance

# Filters and Weights



# Some Common Filters

---

## ➤ **AvailabilityZoneFilter**

- Return hosts where node\_availability\_zone name is the same as the one requested

## ➤ **RamFilter**

- Return hosts where  $(\text{free\_ram} * \text{ram\_allocation\_ration})$  is greater than requested ram.

## ➤ **ComputerFilter**

- Return hosts where asked instance\_type (with extra\_specs) match capabilities

# Some Common Filters (cont.)

---

## ➤ **DiskFilter**

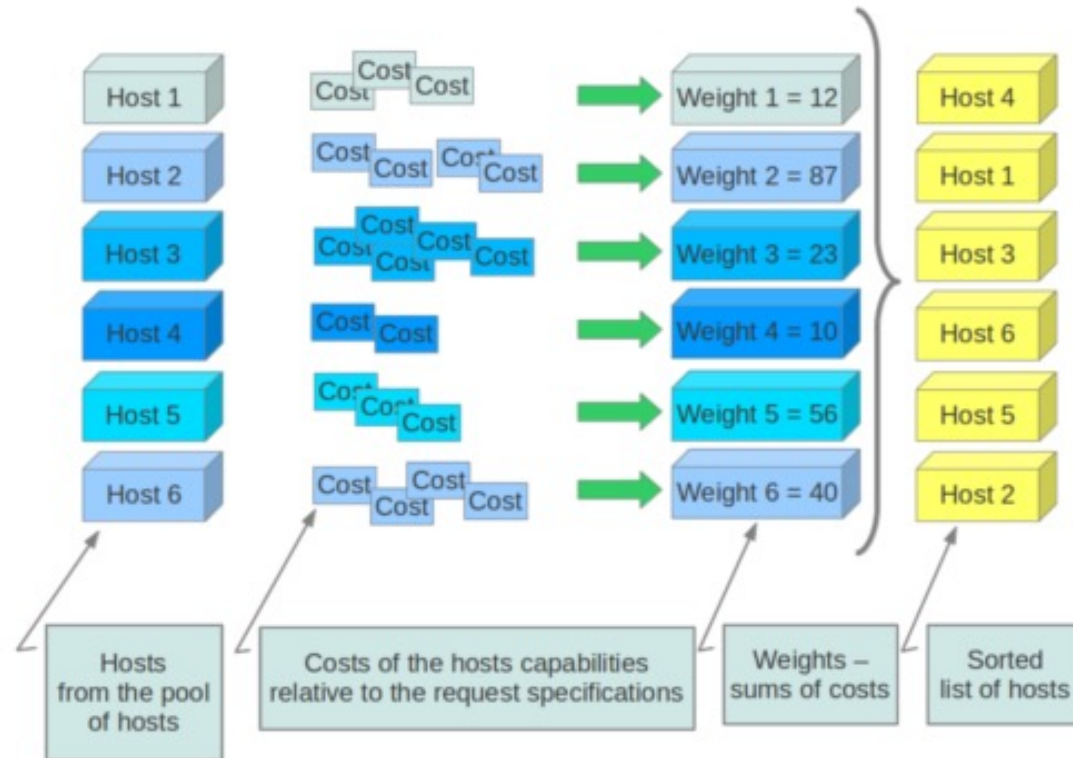
- Returns hosts with sufficient disk space available for root and ephemeral storage.

## ➤ **RetryFilter**

- Filters out hosts that have already been attempted for scheduling purposes.

# Weights

- Scheduler applies cost function on each host & calculates the weight.



<https://docs.openstack.org/nova/latest/admin/scheduling.html>



# Some Possible Cost Functions

---

- Considering free RAM among filtered hosts.
  - Highest free RAM wins.
- Considering least workload (e.g., IO ops) among filtered hosts.
- Can consider any specific metric we want to consider in a similar fashion.
  - Can be enabled from configuration file.

```
weight = w1_multiplier * norm(w1) + w2_multiplier * norm(w2) + ...
```

# Weights (cont.)

---

## ➤ **RAMWeigher**

- Compute weight based on available RAM on the compute node.  
Sort with the largest weight winning.

## ➤ **CPUWeigher**

- Compute weight based on available vCPUs on the compute node.  
Sort with the largest weight winning.

## ➤ **DiskWeigher**

- Hosts are weighted and sorted by free disk space with the largest weight winning.

# Weights (cont.)

---

## ➤ MetricWeigher

- This weigher can compute the weight based on the compute node host's various metrics.
- The to-be weighed metrics and their weighing ration are specified in the configuration file as the followings:

```
[metrics]  
weight_setting = name1=1.0, name2=-1.0
```

# General Cost Function

---

```
weight = w1_multiplier * norm(w1) + w2_multiplier * norm(w2) + ...
```

Metric	Range
CPU utilization	(0, 100) usage percentage
Outbound network traffic	(0, 10 <sup>9</sup> ) byte per second

# *Least* Loaded Server with No Normalization

---

Weight (Load) = 1 \* (CPU utilization) + 1\* (Outbound network traffic )

	CPU utilization	Outbound network traffic
Host1	95	100000
Host2	10	100090

# *Least* Loaded Server Without Normalization

---

Weight (Load) =  $1 * (\text{CPU utilization}) + 1 * (\text{Outbound network traffic})$

	Weight
<b>Host1</b>	$(1 * 95) + (1 * 100000) = 100095$ ✓
Host2	$(1 * 10) + (1 * 100090) = 100100$

Host1 is selected!

Not good 😞

# Min-Max Normalization

---

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

As we discussed in the class, the min and max come from the nature of data

# Getting Back to the Previous Example

---

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

	CPU utilization	Outbound network traffic
Host1	95	100000
Host2	10	100090



	CPU utilization	Outbound network traffic
Host1	$(95-0)/(100-0)=0.95$	$(100000-0)/(10^9-0)=0.0001$
Host2	$(10-0)/(100-0)=0.1$	$(100090-0)/(10^9-0)=0.00010009$



# *Least* Loaded Server with Normalization

---

$$\text{Weight (Load)} = 1 * \text{norm}(\text{CPU utilization}) + 1 * \text{norm}(\text{Outbound network traffic})$$

	Weight
Host1	$(1 * 0.95) + (1 * 0.0001) = 0.9501$
<b>Host2</b>	$(1 * 0.1) + (1 * 0.00010009) = 0.10010009$ ✓

Host2 is selected!

Good job :)

# نمونه سوال امتحانی

## ۵ سوال OpenStack (۲۰ نمره)

الف) به انتخاب خودتان، پنج مولفه اصلی openstack را نام ببرید و وظیفه هر کدام را کوتاه بنویسید. (۱۰ نمره)

ب) فرض کنید openstack رو طوری پیکربندی کردید که برای انتخاب میزبان یک ماشین مجازی از فیلتر RamFilter استفاده نموده و برای وزن دهی، از ترکیب دو شاخص درصد استفاده از CPU و مقدار ترافیک خروجی شبکه میزبان (برحسب کیلو بایت) استفاده می کند. در این پیکربندی ram\_allocation\_ratio برابر با ۱.۵ قرار داده شده و برای شاخص های درصد استفاده از CPU و مقدار ترافیک خروجی وزنی برابر با ۰.۵ در نظر گرفته شده است.

Host name	Zone	Free Memory (GB)	CPU usage (%)	Outbound network traffics (KB)
Server1	A	32	98	1000
Server2	A	8	30	5000
Server3	A	12.5	40	1200
Server4	B	10	80	1030
Server5	B	8	20	500
Server6	B	16	10	1100

درخواستی برای ایجاد ماشین مجازی با 16GB حافظه اصلی و ۴ هسته CPU در Zone B اماده است. مشخص کنید از بین میزبان های جدول بالا کدام میزبان ها از فیلتر عبور می کنند و وزن هر کدام چقدر خواهد بود و کدام یک برای ایجاد ماشین مجازی انتخاب می شود. (۱۰ نمره)

# نکته مهم در مورد پاسخ سوال قبلی

➤ همانطور که در کلاس دیدیم، برای نرمال سازی min-max و تعیین مینیمم و ماکزیمم، بایستی به بازه ممکن برای هر ستون رجوع کرد و مثلاً استفاده از ۵۰۰ به عنوان مینیمم ستون ترافیک خروجی شبکه و ۵۰۰۰ به عنوان ماکزیمم این ستون، به پاسخ صحیحی منتج نمی شود. در اینجا بایستی مینیمم و ماکزیمم برای این ستون به شکل بازه داده می شد. به عنوان مثال باید بازه  $[0, 125000\text{KB}]$  داده می شد و شما مینیمم این ستون را 0 و ماکزیمم آن را 125000 در نظر می گرفتید. از طرفی برای درصد استفاده از پردازنده، بایستی بازه  $[0, 100]$  داده می شد. از طرفی Zone B در این سوال، نکته انحرافی است. این تیپ سوال رو جدی بگیرید (: